

A Flow-Sensitive Refinement Type System for Verifying eBPF Programs



CERTORA



Ameer Hamza¹, **Lucas Zavalía**¹, Arie Gurfinkel², Jorge A. Navas³,
Grigory Fedyukovich¹

¹Florida State University

²University of Waterloo

³Certora, Inc.

What is eBPF?

What is eBPF?

- System for hotloading code into kernel space

What is eBPF?

- System for hotloading code into kernel space
- Commonly Used for

What is eBPF?

- System for hotloading code into kernel space
- Commonly Used for
 - Packet filtering

What is eBPF?

- System for hotloading code into kernel space
- Commonly Used for
 - Packet filtering
 - Process monitoring

What is eBPF?

- System for hotloading code into kernel space
- Commonly Used for
 - Packet filtering
 - Process monitoring
 - Security auditing

What is eBPF?

- System for hotloading code into kernel space
- Commonly Used for
 - Packet filtering
 - Process monitoring
 - Security auditing
 - etc.

What is eBPF?

- System for hotloading code into kernel space
- Commonly Used for
 - Packet filtering
 - Process monitoring
 - Security auditing
 - etc.
- Used throughout industry

What is eBPF?

- System for hotloading code into kernel space
- Commonly Used for
 - Packet filtering
 - Process monitoring
 - Security auditing
 - etc.
- Used throughout industry



What is eBPF?

- Hotloading kernel extensions is unsafe in general

What is eBPF?

- Hotloading kernel extensions is unsafe in general
 - May reveal secrets

What is eBPF?

- Hotloading kernel extensions is unsafe in general
 - May reveal secrets
 - May crash the whole OS

What is eBPF?

- Hotloading kernel extensions is unsafe in general
 - May reveal secrets
 - May crash the whole OS
 - etc.

What is eBPF?

- Hotloading kernel extensions is unsafe in general
 - May reveal secrets
 - May crash the whole OS
 - etc.
- All eBPF programs must be “formally verified”

What is eBPF?

- Hotloading kernel extensions is unsafe in general
 - May reveal secrets
 - May crash the whole OS
 - etc.
- All eBPF programs must be “formally verified”
- eBPF programs are restricted

What is eBPF?

- Hotloading kernel extensions is unsafe in general
 - May reveal secrets
 - May crash the whole OS
 - etc.
- All eBPF programs must be “formally verified”
- eBPF programs are restricted

User Space

Kernel Space

What is eBPF?

- Hotloading kernel extensions is unsafe in general
 - May reveal secrets
 - May crash the whole OS
 - etc.
- All eBPF programs must be “formally verified”
- eBPF programs are restricted

User Space

eBPF Program

Kernel Space

What is eBPF?

- Hotloading kernel extensions is unsafe in general
 - May reveal secrets
 - May crash the whole OS
 - etc.
- All eBPF programs must be “formally verified”
- eBPF programs are restricted

User Space

eBPF Program

Kernel Space

Verifier

What is eBPF?

- Hotloading kernel extensions is unsafe in general
 - May reveal secrets
 - May crash the whole OS
 - etc.
- All eBPF programs must be “formally verified”
- eBPF programs are restricted

User Space

Kernel Space

eBPF Program

Verifier

Load

What is eBPF?

- Hotloading kernel extensions is unsafe in general
 - May reveal secrets
 - May crash the whole OS
 - etc.
- All eBPF programs must be “formally verified”
- eBPF programs are restricted

User Space

Reject

Kernel Space

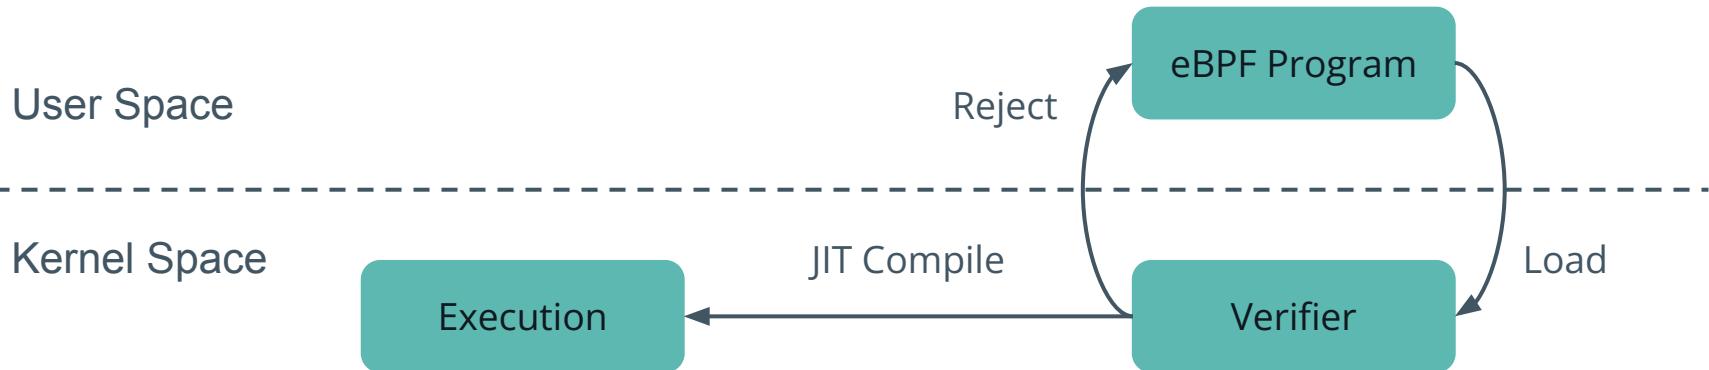
eBPF Program

Verifier

Load

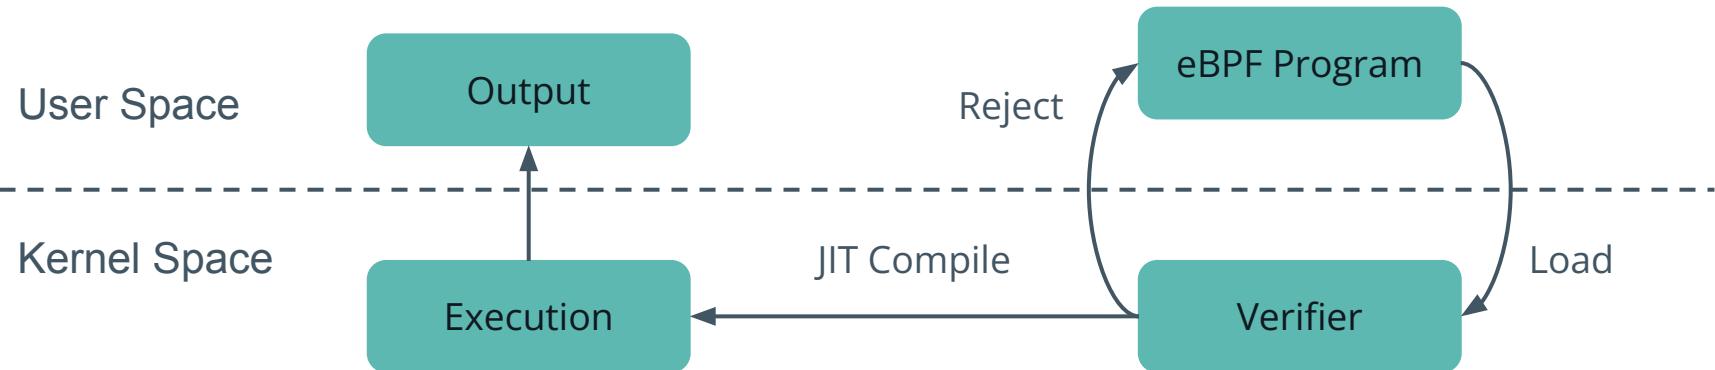
What is eBPF?

- Hotloading kernel extensions is unsafe in general
 - May reveal secrets
 - May crash the whole OS
 - etc.
- All eBPF programs must be “formally verified”
- eBPF programs are restricted



What is eBPF?

- Hotloading kernel extensions is unsafe in general
 - May reveal secrets
 - May crash the whole OS
 - etc.
- All eBPF programs must be “formally verified”
- eBPF programs are restricted



- **Information Flow Safety:** An eBPF program never exposes the address space of the kernel

- **Information Flow Safety:** An eBPF program never exposes the address space of the kernel
- **Memory Safety:** An eBPF program accesses only certain memory regions

- **Information Flow Safety:** An eBPF program never exposes the address space of the kernel
- **Memory Safety:** An eBPF program accesses only certain memory regions

```
0  int packet_proc(struct xdp_md * ctx) {
1      void * data = (void *)(long)ctx->data;
2      void * data_end = (void *)(long)ctx->data_end;
3      if (data > data_end) exit 1;
4      *(int *) data = *(int *)(data & 0x3A53E170);
5      return 0;
6  }
```

- **Information Flow Safety:** An eBPF program never exposes the address space of the kernel
- **Memory Safety:** An eBPF program accesses only certain memory regions

```
0 int packet_proc(struct xdp_md * ctx) {  
1     void * data = (void *) (long) ctx->data;  
2     void * data_end = (void *) (long) ctx->data_end;  
3     if (data > data_end) exit 1;  
4     *(int *) data = *(int *) (data & 0x3A53E170);  
5     return 0;  
6 }
```

Context is passed to eBPF program from OS

- **Information Flow Safety:** An eBPF program never exposes the address space of the kernel
- **Memory Safety:** An eBPF program accesses only certain memory regions

```
0  int packet_proc(struct xdp_md * ctx) {  
1      void * data = (void *)(long)ctx->data;  
2      void * data_end = (void *)(long)ctx->data_end;  
3      if (data > data_end) exit 1;  
4      *(int *) data = *(int *) (data & 0x3A53E170);  
5      return 0;  
6  }
```

Load pointer to
beginning of packet

- **Information Flow Safety:** An eBPF program never exposes the address space of the kernel
- **Memory Safety:** An eBPF program accesses only certain memory regions

```
0  int packet_proc(struct xdp_md * ctx) {  
1      void * data = (void *) (long) ctx->data;  
2      void * data_end = (void *) (long) ctx->data_end;  
3      if (data > data_end) exit 1;  
4      *(int *) data = *(int *) (data & 0x3A53E170);  
5      return 0;  
6  }
```

Load end of
packet

- **Information Flow Safety:** An eBPF program never exposes the address space of the kernel
- **Memory Safety:** An eBPF program accesses only certain memory regions

```
0  int packet_proc(struct xdp_md * ctx) {  
1      void * data = (void *)(long)ctx->data;  
2      void * data_end = (void *)(long)ctx->data_end;  
3      if (data > data_end) exit 1;  
4      *(int *) data = *(int *) (data + 0x3A53E170);  
5      return 0;  
6  }
```

Exit condition

- **Information Flow Safety:** An eBPF program never exposes the address space of the kernel
- **Memory Safety:** An eBPF program accesses only certain memory regions

```
0  int packet_proc(struct xdp_md * ctx) {  
1      void * data = (void *)(long)ctx->data;  
2      void * data_end = (void *)(long)ctx->data_end;  
3      if (data > data_end) exit 1;  
4      *(int *) data = *(int *)(data & 0x3A53E170);  
5      return 0;  
6  }
```

Load/Store 4 bytes

- **Information Flow Safety:** An eBPF program never exposes the address space of the kernel
- **Memory Safety:** An eBPF program accesses only certain memory regions

```
0  int packet_proc(struct xdp_md * ctx) {  
1      void * data = (void *) (long)ctx->data;  
2      void * data_end = (void *) (long)ctx->data_end;  
3      if (data > data_end) exit 1;  
4      *(int *) data = *(int *) (data & 0x3A53E170);  
5      return 0;  
6  }
```

UNSAFE

- **Information Flow Safety:** An eBPF program never exposes the address space of the kernel
- **Memory Safety:** An eBPF program accesses only certain memory regions

```
0  int packet_proc(struct xdp_md * ctx) {  
1      void * data = (void *)(long)ctx->data;  
2      void * data_end = (void *)(long)ctx->data_end;  
3      if (data > data_end) exit 1;  
4      *(int *) data = *(int *)(data & 0x3A53E170);  
5      return 0;  
6  }
```

UNSAFE

- **Information Flow Safety:** An eBPF program never exposes the address space of the kernel
- **Memory Safety:** An eBPF program accesses only certain memory regions

```
0  int packet_proc(struct xdp_md * ctx) {
1      void * data = (void *)(long)ctx->data;
2      void * data_end = (void *)(long)ctx->data_end;
3      if (data + sizeof(int) > data_end) exit 1;
4      *(int *) data = *(int *)(data & 0x3A53E170);
5      return 0;
6  }
```

- **Information Flow Safety:** An eBPF program never exposes the address space of the kernel
- **Memory Safety:** An eBPF program accesses only certain memory regions

```
0  int packet_proc(struct xdp_md * ctx) {  
1      void * data = (void *) (long) ctx->data;  
2      void * data_end = (void *) (long) ctx->data_end;  
3      if (data + sizeof(int) > data_end) exit 1;  
4      *(int *) data = *(int *) (data & 0x3A53E170);  
5      return 0;  
6  }
```

SAFE ✓


```
1 []
2 Pre-invariant : [
3     meta_offset=[ -4098, 0 ],
4     packet_size=[ 0, 65534 ],
5     r1.ctx_offset=0, r1.svalue=[ 1, 2147418112 ], r1.type=ctx,
6     r10.stack_offset=512, r10.svalue=[ 512, 2147418112 ], r10.type=stack]
7 Stack: Numbers -> {}
8 entry:
9   r0 = 1;
10  assert r1.type in {ctx, stack, packet, shared};
11  assert valid_access(r1.offset, width=4) for read;
12  r2 = *(u32 *)(r1 + 0);
13  assert r1.type in {ctx, stack, packet, shared};
14  assert valid_access(r1.offset+4, width=4) for read;
15  r1 = *(u32 *)(r1 + 4);
16  assert valid_access(r2.offset) for comparison/subtraction;
17  assert valid_access(r1.offset) for comparison/subtraction;
18  assert r2.type in {number, ctx, stack, packet, shared};
19  assert r2.type == r1.type in {ctx, stack, packet};
20  goto 3:4,3:8;
```

```
126 3:4: Upper bound must be at most packet_size (valid_access(r2.offset, width=4) for read)
127 3:4: Upper bound must be at most packet_size (valid_access(r2.offset, width=4) for write)
128 3:8: Code is unreachable after 3:8
129
130 0,0.000811,5504
```

Existing Verifiers

Existing Verifiers

- **Linux Verifier:**

Existing Verifiers

- **Linux Verifier:**
 - Uses symbolic execution

Existing Verifiers

- **Linux Verifier:**
 - Uses symbolic execution
 - First eBPF verifier

Existing Verifiers

- **LinuxVerifier:**
 - Uses symbolic execution
 - First eBPF verifier
- **PREVAIL [PLDI, 2019]:**

Existing Verifiers

- **LinuxVerifier:**
 - Uses symbolic execution
 - First eBPF verifier
- **PREVAIL [PLDI, 2019]:**
 - Uses abstract interpretation

Existing Verifiers

- **LinuxVerifier:**
 - Uses symbolic execution
 - First eBPF verifier
- **PREVAIL [PLDI, 2019]:**
 - Uses abstract interpretation
 - Part of eBPF for Windows

Existing Verifiers

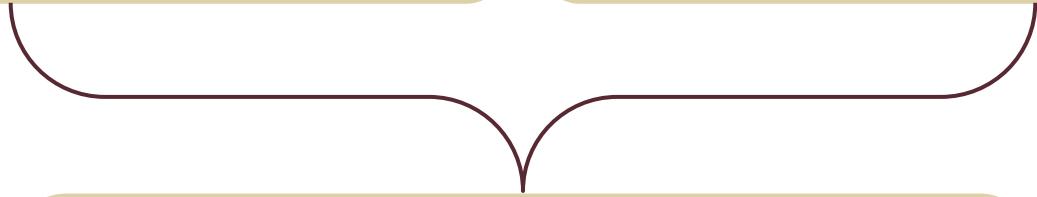
- **Linux Verifier:**
 - Uses symbolic execution
 - First eBPF verifier
- **PREVAIL [PLDI, 2019]:**
 - Uses abstract interpretation
 - Part of eBPF for Windows

- **Different Verifier Same Problems:**

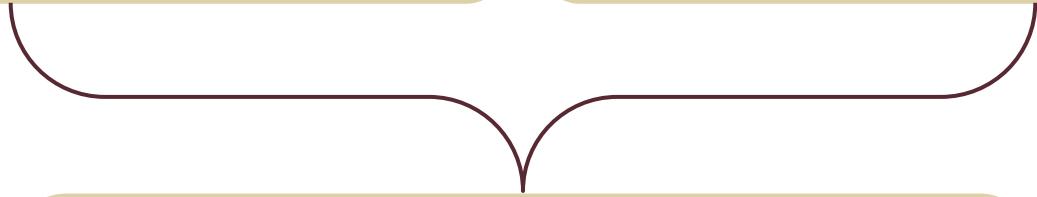
Existing Verifiers

- **Linux Verifier:**
 - Uses symbolic execution
 - First eBPF verifier
 - **PREVAIL [PLDI, 2019]:**
 - Uses abstract interpretation
 - Part of eBPF for Windows
-
- **DifferentVerifierSameProblems:**
 - Verifier is monolithic

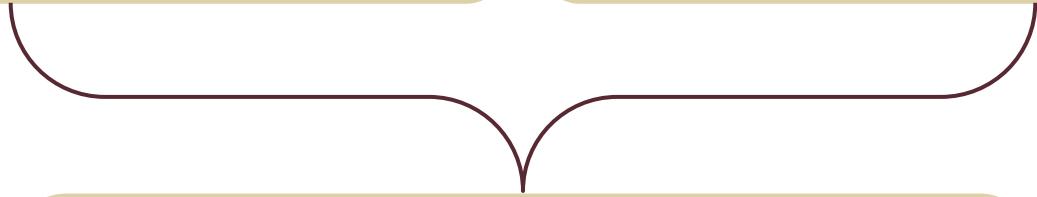
Existing Verifiers

- **Linux Verifier:**
 - Uses symbolic execution
 - First eBPF verifier
 - **PREVAIL [PLDI, 2019]:**
 - Uses abstract interpretation
 - Part of eBPF for Windows
- 
- **DifferentVerifierSameProblems:**
 - Verifier is monolithic
 - Lives entirely in kernel space

Existing Verifiers

- **LinuxVerifier:**
 - Uses symbolic execution
 - First eBPF verifier
 - **PREVAIL [PLDI, 2019]:**
 - Uses abstract interpretation
 - Part of eBPF for Windows
- 
- **Different Verifier Same Problems:**
 - Verifier is monolithic
 - Lives entirely in kernel space
 - Users cannot provide guidance

Existing Verifiers

- **LinuxVerifier:**
 - Uses symbolic execution
 - First eBPF verifier
 - **PREVAIL [PLDI, 2019]:**
 - Uses abstract interpretation
 - Part of eBPF for Windows
- 
- **Different Verifier Same Problems:**
 - Verifier is monolithic
 - Lives entirely in kernel space
 - Users cannot provide guidance
 - Verifier logs are hard to read

- Type systems addresses architectural problems from existing verifiers

Type Based Verification

- Type systems addresses architectural problems from existing verifiers
 - *Type inference* done in user space

Type Based Verification

- Type systems addresses architectural problems from existing verifiers
 - *Type inference* done in user space
 - *Type checking* done in kernel space

Type Based Verification

- Type systems addresses architectural problems from existing verifiers
 - *Type inference* done in user space
 - *Type checking* done in kernel space
- Types are relatively more intuitive for programmers

Type Based Verification

- Type systems addresses architectural problems from existing verifiers
 - *Type inference* done in user space
 - *Type checking* done in kernel space
- Types are relatively more intuitive for programmers

User Space

Kernel Space

Type Based Verification

- Type systems addresses architectural problems from existing verifiers
 - *Type inference* done in user space
 - *Type checking* done in kernel space
- Types are relatively more intuitive for programmers

User Space

eBPF Program

Kernel Space

Type Based Verification

- Type systems addresses architectural problems from existing verifiers
 - *Type inference* done in user space
 - *Type checking* done in kernel space
- Types are relatively more intuitive for programmers

User Space



Kernel Space

Type Based Verification

- Type systems addresses architectural problems from existing verifiers
 - *Type inference* done in user space
 - *Type checking* done in kernel space
- Types are relatively more intuitive for programmers

User Space



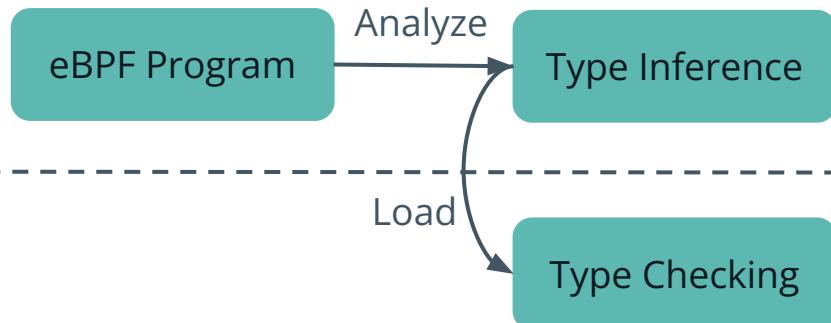
Kernel Space

Type Checking

Type Based Verification

- Type systems addresses architectural problems from existing verifiers
 - *Type inference* done in user space
 - *Type checking* done in kernel space
- Types are relatively more intuitive for programmers

User Space



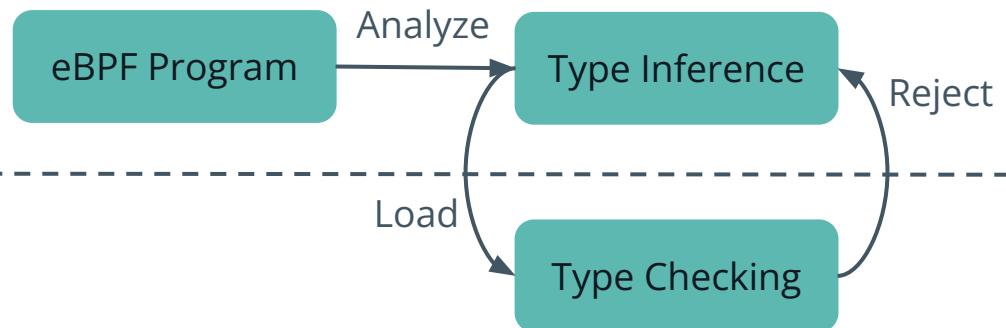
Kernel Space

Type Based Verification

- Type systems addresses architectural problems from existing verifiers
 - *Type inference* done in user space
 - *Type checking* done in kernel space
- Types are relatively more intuitive for programmers

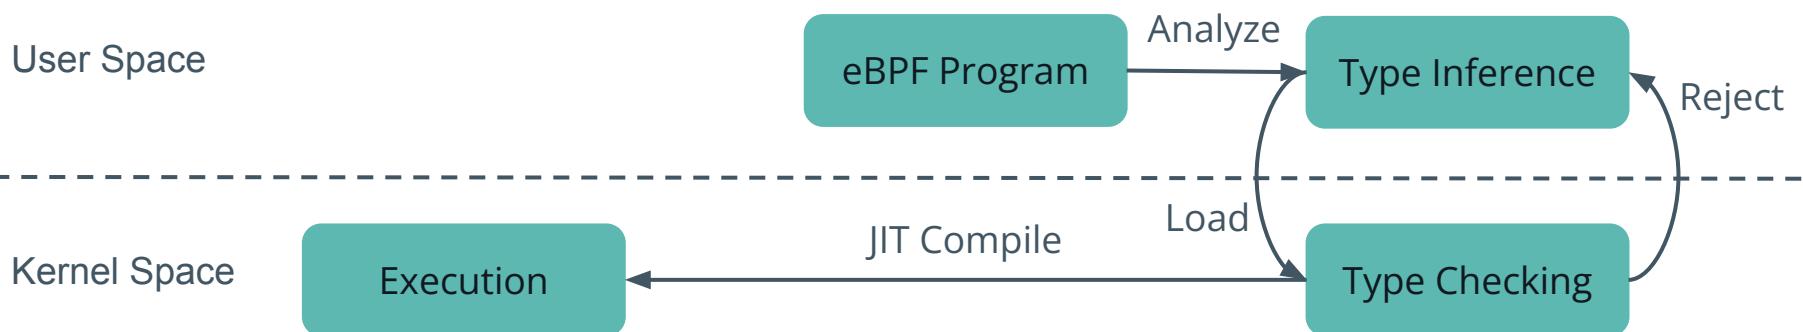
User Space

Kernel Space



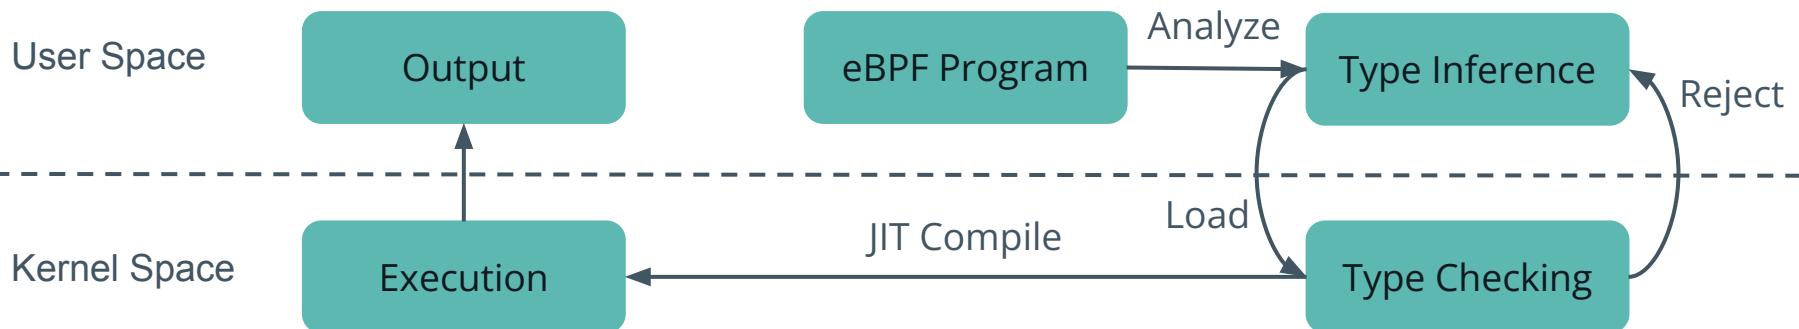
Type Based Verification

- Type systems addresses architectural problems from existing verifiers
 - *Type inference* done in user space
 - *Type checking* done in kernel space
- Types are relatively more intuitive for programmers



Type Based Verification

- Type systems addresses architectural problems from existing verifiers
 - *Type inference* done in user space
 - *Type checking* done in kernel space
- Types are relatively more intuitive for programmers



Basic Types

Three Basic Kinds Objects:

Basic Types

Three Basic Kinds Objects:

- Pointers

Basic Types

Three Basic Kinds Objects:

- Pointers
- Numbers

Basic Types

Three Basic Kinds Objects:

- Pointers
- Numbers
- File Descriptors

Basic Types

Three Basic Kinds Objects:

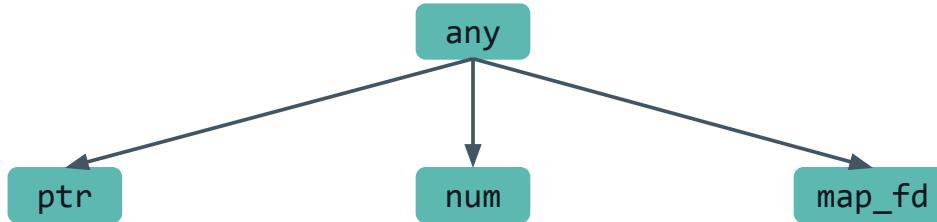
- Pointers
- Numbers
- File Descriptors

any

Basic Types

Three Basic Kinds Objects:

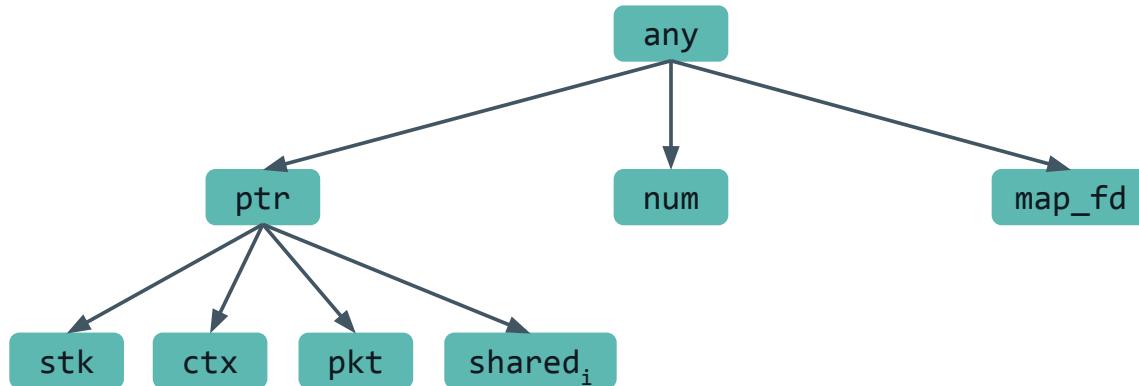
- Pointers
- Numbers
- File Descriptors



Basic Types

Three Basic Kinds Objects:

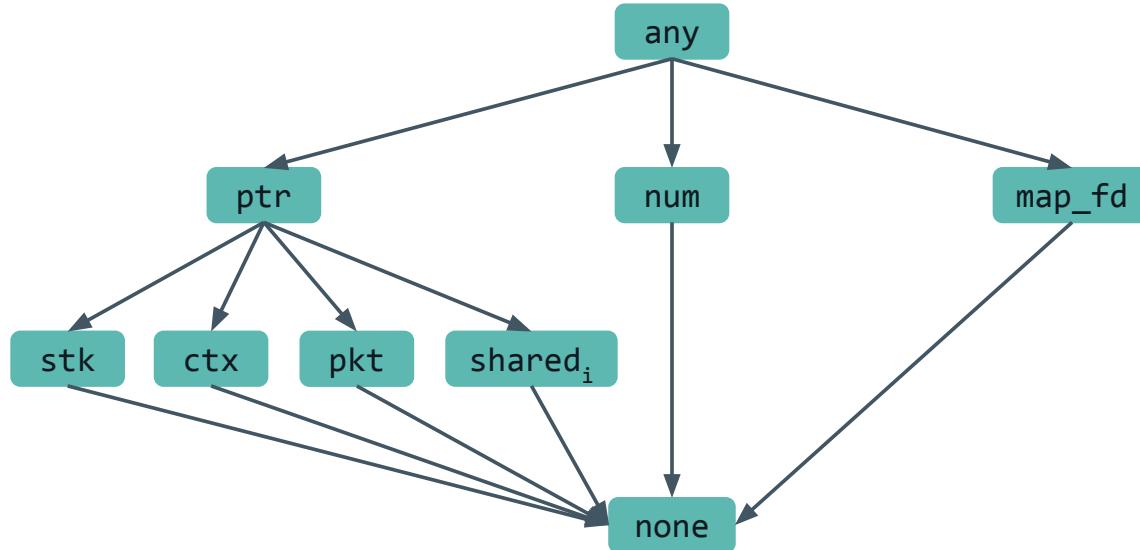
- Pointers
- Numbers
- File Descriptors



Basic Types

Three Basic Kinds Objects:

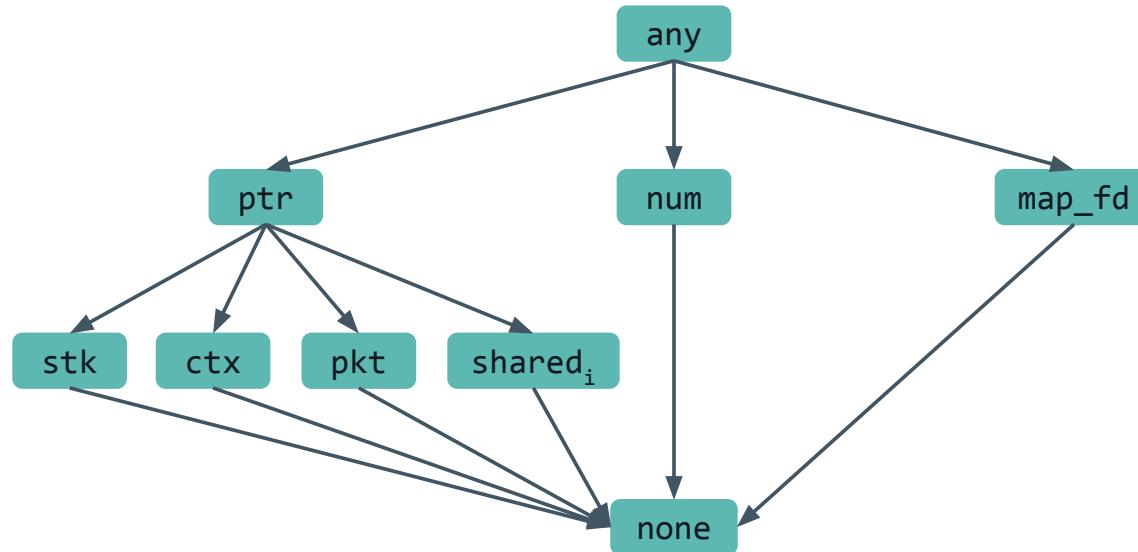
- Pointers
- Numbers
- File Descriptors



Three Basic Kinds Objects:

- Pointers
- Numbers
- File Descriptors

Information-Flow Safety:

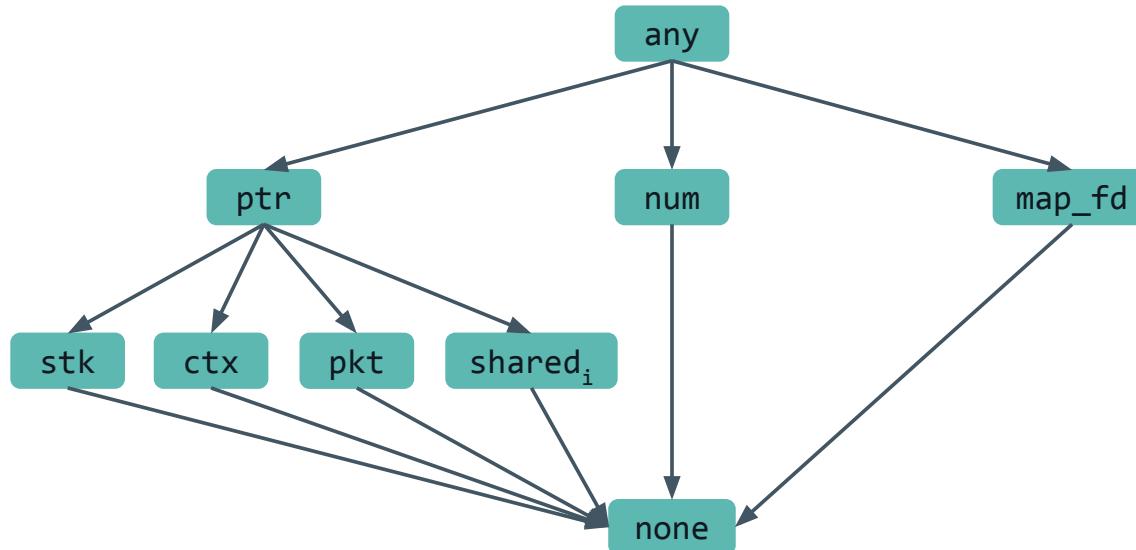


Three Basic Kinds Objects:

- Pointers
- Numbers
- File Descriptors

Information-Flow Safety:

- Cannot store pointer in `sharedi` or `pkt` regions

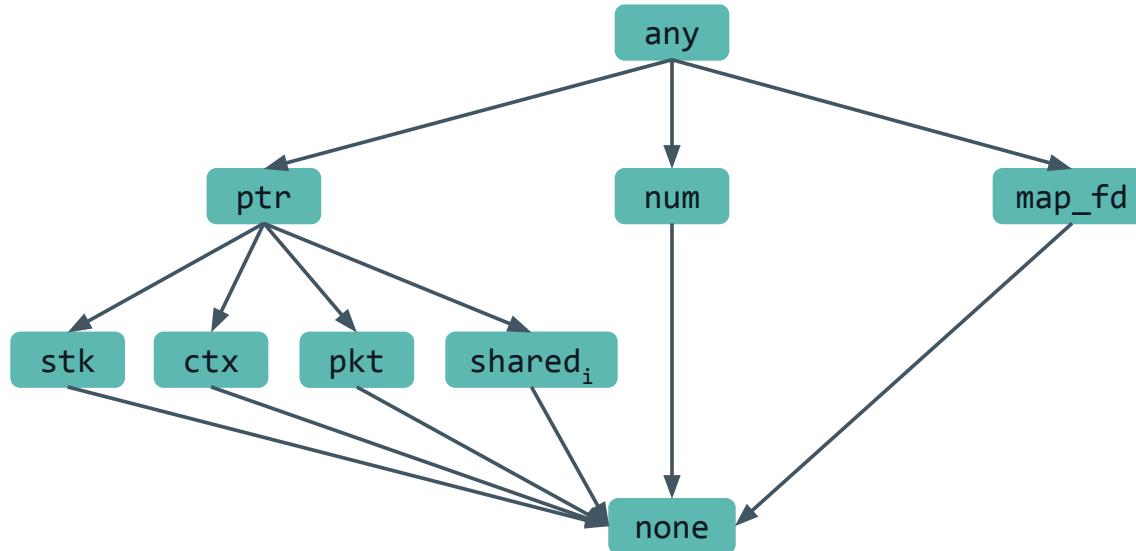


Three Basic Kinds Objects:

- Pointers
- Numbers
- File Descriptors

Information-Flow Safety:

- Cannot store pointer in `sharedi` or `pkt` regions
- OK for `stk` region



Type Environments:

Type Environments:

- Γ = Mapping of registers to types

Flow Sensitivity and Type Refinement

Type Environments:

- Γ = Mapping of registers to types
- Δ = Abstraction of memory regions

Flow Sensitivity and Type Refinement

Type Environments:

- Γ = Mapping of registers to types
- Δ = Abstraction of memory regions

Type Judgements:

Type Environments:

- Γ = Mapping of registers to types
- Δ = Abstraction of memory regions

Type Judgements:

- $\Gamma, \Delta \vdash E : \tau$

Type Environments:

- Γ = Mapping of registers to types
- Δ = Abstraction of memory regions

Type Judgements:

- $\Gamma, \Delta \vdash E : \tau$
- $\Gamma, \Delta \vdash i \dashv \Gamma', \Delta'$

Type Environments:

- Γ = Mapping of registers to types
- Δ = Abstraction of memory regions

Type Judgements:

- $\Gamma, \Delta \vdash E : \tau$
- $\Gamma, \Delta \vdash i \dashv \Gamma', \Delta'$

$$\{\nu:\tau \mid P(\nu)\}$$

Type Environments:

- Γ = Mapping of registers to types
- Δ = Abstraction of memory regions

Type Judgements:

- $\Gamma, \Delta \vdash E : \tau$
- $\Gamma, \Delta \vdash \iota \dashv \Gamma', \Delta'$

$$\{v:\tau \mid P(v)\}$$

Flow Sensitivity and Type Refinement

Type Environments:

- Γ = Mapping of registers to types
- Δ = Abstraction of memory regions

Type Judgements:

- $\Gamma, \Delta \vdash E : \tau$
- $\Gamma, \Delta \vdash i \dashv \Gamma', \Delta'$

$$\{v:\tau \mid P(v)\}$$

Value Variable

Flow Sensitivity and Type Refinement

Type Environments:

- Γ = Mapping of registers to types
- Δ = Abstraction of memory regions

Type Judgements:

- $\Gamma, \Delta \vdash E : \tau$
- $\Gamma, \Delta \vdash i \dashv \Gamma', \Delta'$

$$\{v:\tau \mid P(v)\}$$

Value Variable

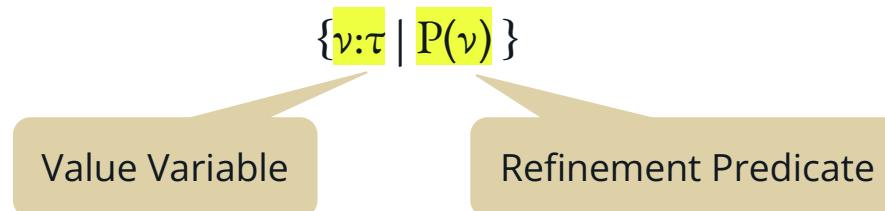
Flow Sensitivity and Type Refinement

Type Environments:

- Γ = Mapping of registers to types
- Δ = Abstraction of memory regions

Type Judgements:

- $\Gamma, \Delta \vdash E : \tau$
- $\Gamma, \Delta \vdash \iota \dashv \Gamma', \Delta'$



Flow Sensitivity and Type Refinement

Type Environments:

- Γ = Mapping of registers to types
- Δ = Abstraction of memory regions

Type Judgements:

- $\Gamma, \Delta \vdash E : \tau$
- $\Gamma, \Delta \vdash \iota \dashv \Gamma', \Delta'$

$\{v:\tau \mid P(v)\}$

Value Variable

Refinement Predicate

$\Gamma, \Delta \vdash r1 := 12$

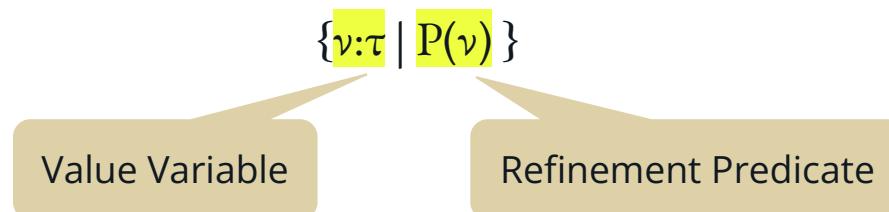
Flow Sensitivity and Type Refinement

Type Environments:

- Γ = Mapping of registers to types
- Δ = Abstraction of memory regions

Type Judgements:

- $\Gamma, \Delta \vdash E : \tau$
- $\Gamma, \Delta \vdash \iota \dashv \Gamma', \Delta'$


$$\Gamma, \Delta \vdash r1 := 12$$
$$\dashv \Gamma[r1 \mapsto \{v:\text{num} \mid v = 12\}], \Delta$$

Flow Sensitivity and Type Refinement

Type Environments:

- Γ = Mapping of registers to types
- Δ = Abstraction of memory regions

Type Judgements:

- $\Gamma, \Delta \vdash E : \tau$
- $\Gamma, \Delta \vdash \iota \dashv \Gamma', \Delta'$

$$\{\nu:\tau \mid P(\nu)\}$$

Value Variable

Refinement Predicate

$$\Gamma, \Delta \vdash r1 := 12$$
$$\dashv \Gamma[r1 \mapsto \{\nu:\text{num} \mid \nu = 12\}], \Delta$$

Flow Sensitivity and Type Refinement

Type Environments:

- Γ = Mapping of registers to types
- Δ = Abstraction of memory regions

Type Judgements:

- $\Gamma, \Delta \vdash E : \tau$
- $\Gamma, \Delta \vdash \iota \dashv \Gamma', \Delta'$

$$\{\nu:\tau \mid P(\nu)\}$$

Value Variable

Refinement Predicate

$$\Gamma, \Delta \vdash r1 := 12$$
$$\dashv \Gamma[r1 \mapsto \{\nu:\text{num} \mid \nu = 12\}], \Delta$$
$$\Gamma[r10 \mapsto \{\nu:\text{stk} \mid \nu = 512\}], \Delta \vdash * (r10 - 4) :=_4 12$$

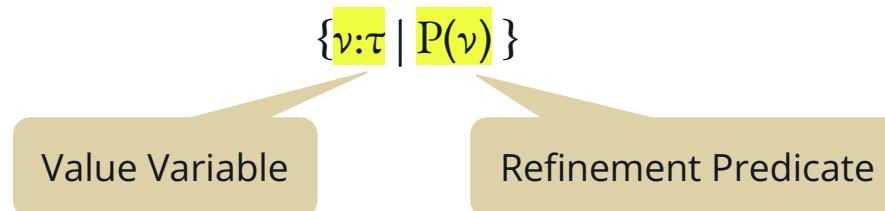
Flow Sensitivity and Type Refinement

Type Environments:

- Γ = Mapping of registers to types
- Δ = Abstraction of memory regions

Type Judgements:

- $\Gamma, \Delta \vdash E : \tau$
- $\Gamma, \Delta \vdash \iota \dashv \Gamma', \Delta'$


$$\Gamma, \Delta \vdash r1 := 12$$
$$\dashv \Gamma[r1 \mapsto \{v:\text{num} \mid v = 12\}], \Delta$$
$$\Gamma[r10 \mapsto \{v:\text{stk} \mid v = 512\}], \Delta \vdash * (r10 - 4) :=_4 12 \dashv \Gamma, \Delta[\text{stk}, 508, \{v:\text{num} \mid v = 12\}]$$

Handling Packet Pointers

- Size of packet unknown at compile time

Handling Packet Pointers

- Size of packet unknown at compile time
- Make sure program has runtime check

Handling Packet Pointers

- Size of packet unknown at compile time
- Make sure program has runtime check
- Special symbols for the packet (`meta`, `begin`, `end`)

Handling Packet Pointers

- Size of packet unknown at compile time
- Make sure program has runtime check
- Special symbols for the packet (`meta`, `begin`, `end`)
- Invariant: $\text{meta} \leq \text{begin} \leq \text{end}$

Handling Packet Pointers

- Size of packet unknown at compile time
- Make sure program has runtime check
- Special symbols for the packet (`meta`, `begin`, `end`)
- Invariant: `meta ≤ begin ≤ end`
- Slack variables are an optimization

Handling Packet Pointers

- Size of packet unknown at compile time
- Make sure program has runtime check
- Special symbols for the packet (`meta`, `begin`, `end`)
- Invariant: `meta ≤ begin ≤ end`
- Slack variables are an optimization

$$\{\nu:\text{pkt} \mid \nu = \text{begin} + s_0 \wedge s_0 \in [0,4] \wedge \text{begin} + 8 \leq \text{end}\}$$

Handling Packet Pointers

- Size of packet unknown at compile time
- Make sure program has runtime check
- Special symbols for the packet (`meta`, `begin`, `end`)
- Invariant: `meta ≤ begin ≤ end`
- Slack variables are an optimization

$$\{\nu:\text{pkt} \mid \nu = \text{begin} + s_0 \wedge s_0 \in [0,4] \wedge \text{begin} + 8 \leq \text{end}\}$$

The location
where the pointer
is pointing to

Handling Packet Pointers

- Size of packet unknown at compile time
- Make sure program has runtime check
- Special symbols for the packet (`meta`, `begin`, `end`)
- Invariant: `meta ≤ begin ≤ end`
- Slack variables are an optimization

$$\{\nu:\text{pkt} \mid \nu = \text{begin} + s_0 \wedge s_0 \in [0,4] \wedge \text{begin} + 8 \leq \text{end}\}$$

The location
where the pointer
is pointing to

Constraint on
slack variable

Handling Packet Pointers

- Size of packet unknown at compile time
- Make sure program has runtime check
- Special symbols for the packet (`meta`, `begin`, `end`)
- Invariant: `meta ≤ begin ≤ end`
- Slack variables are an optimization

$$\{\nu:\text{pkt} \mid \nu = \text{begin} + s_0 \wedge s_0 \in [0,4] \wedge \text{begin} + 8 \leq \text{end}\}$$

The location
where the pointer
is pointing to

Constraint on
slack variable

Safety Condition

Handling Joins

Handling Joins

- Branches can create diverging typings

Handling Joins

- Branches can create diverging typings
- What should the types be at the join point?

Handling Joins

- Branches can create diverging typings
- What should the types be at the join point?
- Preserve information in common

Handling Joins

- Branches can create diverging typings
- What should the types be at the join point?
- Preserve information in common

$$r0 : \{v:\text{num} \mid v = 12\}$$

Handling Joins

- Branches can create diverging typings
- What should the types be at the join point?
- Preserve information in common

$$r0 : \{v:\text{num} \mid v = 12\}$$
$$r0 : \{v:\text{num} \mid v = 4\}$$

Handling Joins

- Branches can create diverging typings
- What should the types be at the join point?
- Preserve information in common

 $r0 : \{v:\text{num} \mid v = 12\}$ $r0 : \{v:\text{num} \mid v = 4\}$ $r0 : \{v:\text{num} \mid v \in [4, 12]\}$

Handling Joins

- Branches can create diverging typings
- What should the types be at the join point?
- Preserve information in common
- Some typings are incomparable

Handling Joins

- Branches can create diverging typings
- What should the types be at the join point?
- Preserve information in common
- Some typings are incomparable

$$r0 : \{v:\text{num} \mid v = 12\}$$

Handling Joins

- Branches can create diverging typings
- What should the types be at the join point?
- Preserve information in common
- Some typings are incomparable

$$r0 : \{v:\text{num} \mid v = 12\}$$
$$r0 : \{v:\text{stk} \mid v = 512\}$$

Handling Joins

- Branches can create diverging typings
- What should the types be at the join point?
- Preserve information in common
- Some typings are incomparable

 $r\theta : \{v:\text{num} \mid v = 12\}$ $r\theta : \{v:\text{stk} \mid v = 512\}$ $r\theta : \text{Any}$

Typed Example

```
0 int packet_proc(struct xdp_md * ctx) {
1     void * data = (void *)(long)ctx->data;
2     void * data_end = (void *)(long)ctx->data_end;
3     if (data > data_end) exit 1;
4     *(int *) data = *(int *)(data & 0x3A53E170);
5     return 0;
6 }
```

Typed Example

```
0 int packet_proc(struct xdp_md * ctx) {  
1     void * data = (void *)(long)ctx->data;  
2     void * data_end = (void *)(long)ctx->data_end;  
3     if (data > data_end) exit 1;  
4     *(int *) data = *(int *)(data & 0x3A53E170);  
5     return 0;  
6 }
```

data : {v:pkt | v = begin }

Typed Example

```
0 int packet_proc(struct xdp_md * ctx) {  
1     void * data = (void *)(long)ctx->data;  
2     void * data_end = (void *)(long)ctx->data_end;  
3     if (data > data_end) exit 1;  
4     *(int *) data = *(int *)(data & 0x3A53E170);  
5     return 0;  
6 }
```

data_end : {v:pkt | v = end}

Typed Example

```
0 int packet_proc(struct xdp_md * ctx) {  
1     void * data = (void *)(long)ctx->data;  
2     void * data_end = (void *)(long)ctx->data_end;  
3     if (data > data_end) exit 1;  
4     *(int *) data = *(int *)(data & 0x3A53E170);  
5     return 0;  
6 }
```

data : none

Typed Example

```
0 int packet_proc(struct xdp_md * ctx) {  
1     void * data = (void *)(long)ctx->data;  
2     void * data_end = (void *)(long)ctx->data_end;  
3     if (data > data_end) exit 1;  
4     *(int *) data = *(int *)(data & 0x3A53E170);  
5     return 0;  
6 }
```

data : {v:pkt | v = begin \wedge begin \leq end }

Typed Example

```
0 int packet_proc(struct xdp_md * ctx) {
1     void * data = (void *)(long)ctx->data;
2     void * data_end = (void *)(long)ctx->data_end;
3     if (data + sizeof(int) > data_end) exit 1;
4     *(int *) data = *(int *)(data & 0x3A53E170);
5     return 0;
6 }
```

Typed Example

```
0 int packet_proc(struct xdp_md * ctx) {  
1     void * data = (void *)(long)ctx->data;  
2     void * data_end = (void *)(long)ctx->data_end;  
3     if (data + sizeof(int) > data_end) exit 1;  
4     *(int *) data = *(int *)(data & 0x3A53E170);  
5     return 0;  
6 }
```

$$\text{data} : \{v:\text{pkt} \mid v = \text{begin} \wedge \text{begin} + 4 \leq \text{end}\}$$

Traverse control graph:

Traverse control graph:

- For one basic block:

Traverse control graph:

- For one basic block:
 - Join with predecessors (if any)

Traverse control graph:

- For one basic block:
 - Join with predecessors (if any)
 - Infer types for instructions

Traverse control graph:

- For one basic block:
 - Join with predecessors (if any)
 - Infer types for instructions
- For loop heads:

Type Inference

Traverse control graph:

- For one basic block:
 - Join with predecessors (if any)
 - Infer types for instructions
- For loop heads:
 - Infer types for basic blocks

Type Inference

Traverse control graph:

- For one basic block:
 - Join with predecessors (if any)
 - Infer types for instructions
- For loop heads:
 - Infer types for basic blocks
 - Widen until a fixed point is reached

- Type inference algorithm implemented in a tool called VeRefine

- Type inference algorithm implemented in a tool called VeRefine
- Implemented as an abstract domain in the PREVAIL framework

- Type inference algorithm implemented in a tool called VeRefine
- Implemented as an abstract domain in the PREVAIL framework
- Evaluated on 420 industrial and synthetic benchmarks

- Type inference algorithm implemented in a tool called VeRefine
- Implemented as an abstract domain in the PREVAIL framework
- Evaluated on 420 industrial and synthetic benchmarks

- Type inference algorithm implemented in a tool called VeRefine
- Implemented as an abstract domain in the PREVAIL framework
- Evaluated on 420 industrial and synthetic benchmarks
 - Cilium, Suricata, OVS, Falco, etc.

Implementation and Evaluation

- Type inference algorithm implemented in a tool called VeRefine
- Implemented as an abstract domain in the PREVAIL framework
- Evaluated on 420 industrial and synthetic benchmarks
 - Cilium, Suricata, OVS, Falco, etc.

	# benchmarks solved	# benchmarks verified	Unique benchs verified	Average time	Maximum time
VEREFINE	420	406 (337/69)	15	0.06s	1.12s
PREVAIL	401	401 (332/69)	10	0.51s	14.59s

Implementation and Evaluation

- Type inference algorithm implemented in a tool called VeRefine
- Implemented as an abstract domain in the PREVAIL framework
- Evaluated on 420 industrial and synthetic benchmarks
 - Cilium, Suricata, OVS, Falco, etc.
- Less precise than prevail (but still sufficiently precise)

	# benchmarks solved	# benchmarks verified	Unique benchs verified	Average time	Maximum time
VEREFINE	420	406 (337/69)	15	0.06s	1.12s
PREVAIL	401	401 (332/69)	10	0.51s	14.59s

Implementation and Evaluation

- Type inference algorithm implemented in a tool called VeRefine
- Implemented as an abstract domain in the PREVAIL framework
- Evaluated on 420 industrial and synthetic benchmarks
 - Cilium, Suricata, OVS, Falco, etc.
- Less precise than prevail (but still sufficiently precise)
- Faster than PREVAIL

	# benchmarks solved	# benchmarks verified	Unique benchs verified	Average time	Maximum time
VEREFINE	420	406 (337/69)	15	0.06s	1.12s
PREVAIL	401	401 (332/69)	10	0.51s	14.59s

Summary:

Conclusion

Summary:

- Type system for verifying eBPF Programs

Conclusion

Summary:

- Type system for verifying eBPF Programs
- Enhanced debuggability through the use of types

Conclusion

Summary:

- Type system for verifying eBPF Programs
- Enhanced debuggability through the use of types
- Improved modularity of eBPF verification

Conclusion

Summary:

- Type system for verifying eBPF Programs
- Enhanced debuggability through the use of types
- Improved modularity of eBPF verification
- Improved performance

Conclusion

Summary:

- Type system for verifying eBPF Programs
- Enhanced debuggability through the use of types
- Improved modularity of eBPF verification
- Improved performance

Future Work:

Conclusion

Summary:

- Type system for verifying eBPF Programs
- Enhanced debuggability through the use of types
- Improved modularity of eBPF verification
- Improved performance

Future Work:

- Type checker

Conclusion

Summary:

- Type system for verifying eBPF Programs
- Enhanced debuggability through the use of types
- Improved modularity of eBPF verification
- Improved performance

Future Work:

- Type checker
- Termination analysis

Conclusion

Summary:

- Type system for verifying eBPF Programs
- Enhanced debuggability through the use of types
- Improved modularity of eBPF verification
- Improved performance

Future Work:

- Type checker
- Termination analysis



Code Available Here